

ERIK PELLIZZON

Go Developer - Freelancer

BYPASSING THE LINUX NET STACK WITH GO



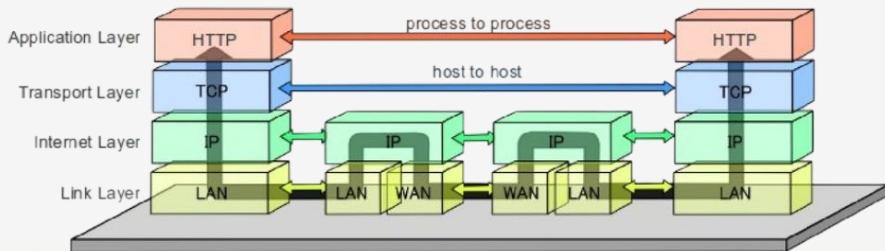
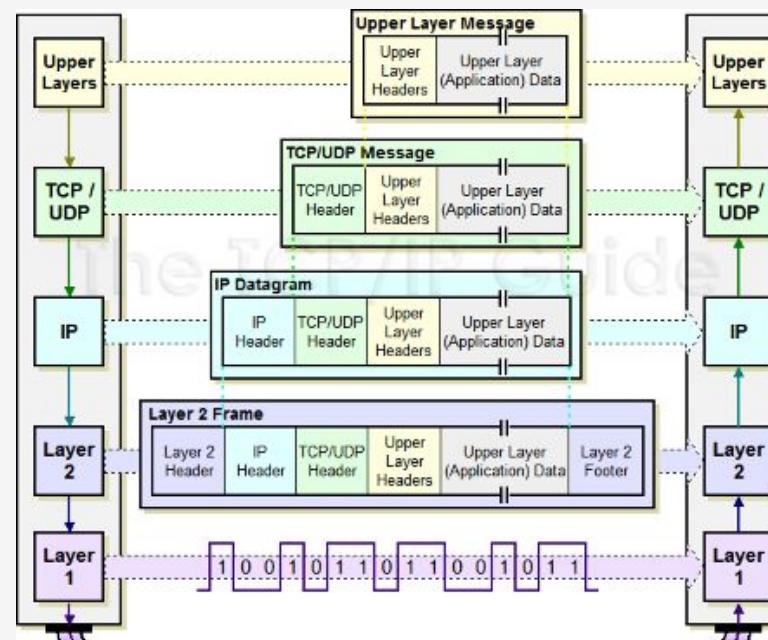
► Introduction

Using XDP and eBPF, we can directly read and write Ethernet frames from the network card and parse them ourselves, gaining complete control over the data and significantly improving performance, basically bypassing all the Linux network layers.

And all this can be done directly from a normal Go application run by the user... How beautiful is it?

Bypassing the Linux net stack with Go

Network

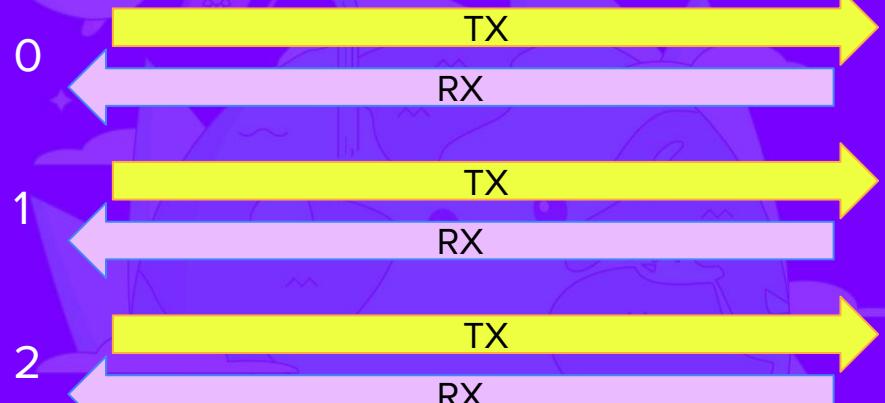


Network

► Network Interface Card (NIC)



Queues



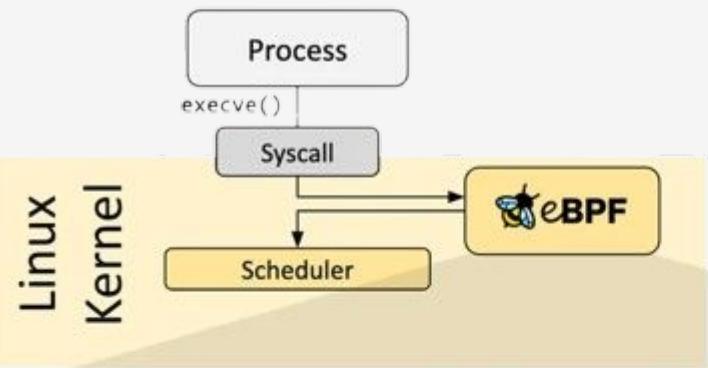
▶ Queues

```
func getInterfaceQueuesNumber(interfaceName string) (int, error) {
    // github.com/vishvananda/netlink
    link, err := netlink.LinkByName(interfaceName)
    if err != nil {
        return 0, err
    }
    queues := min(link.Attrs().NumRxQueues, link.Attrs().NumTxQueues)
    if queues < 1 {
        return 0, errors.New("no queues found")
    }
    return queues, nil
}
```

eBPF



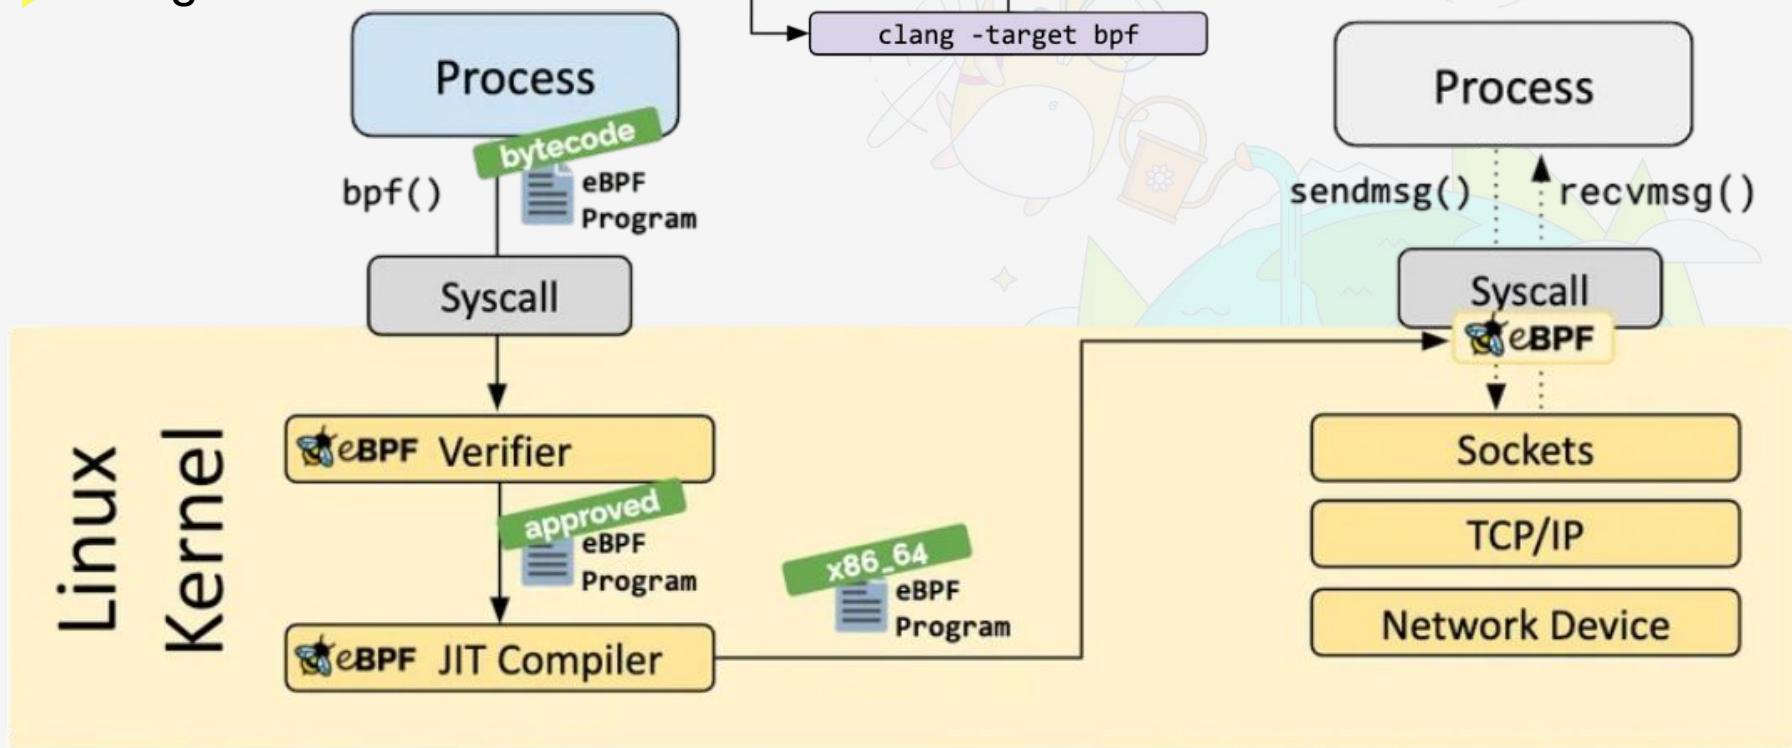
extended Berkeley Packet Filter



- Starting from Linux 3.18 (2014)
- Event-driven (system calls, network events, storage, etc.)
- eBPF program
 - Runs in a VM, within the kernel
 - Bytecode, JIT
 - Must be loaded with **bpf** system call (root permissions required by default)

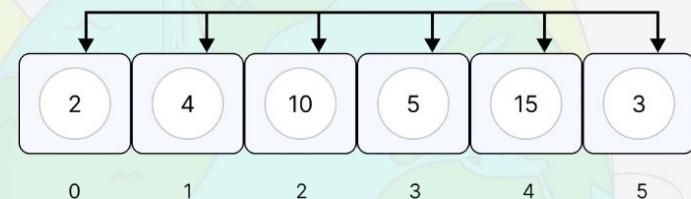
eBPF

Program



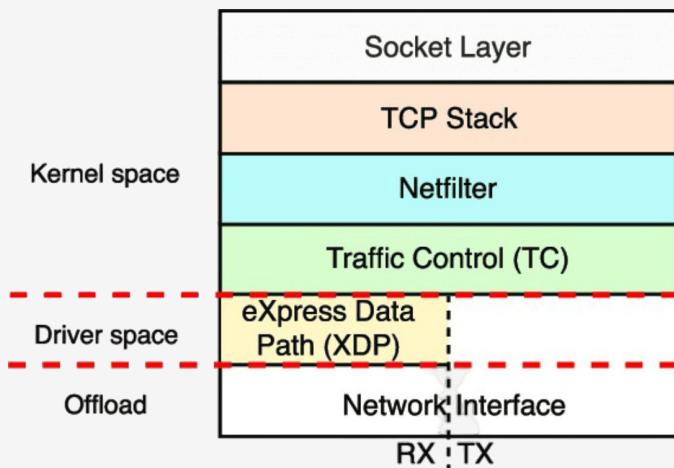
Maps

- They have a type assigned: e.g.
`BPF_MAP_TYPE_HASH` (hashmap) or
`BPF_MAP_TYPE_ARRAY` (array)
- Key-Value data structures
- Used to communicate with other eBPF programs or to receive data from the user space application



XDP


eXpress Data Path




- Starting from Linux 4.8 (2016)
- Built on top of eBPF
- eBPF program can run directly on the NIC (**offload**)
- XDP hook point is inside the NIC driver, called for **every incoming packet**
- When the NIC driver doesn't support XDP, eBPF program can be attached to the traditional net stack (but it's slower)

▶ AF_XDP socket


- AF_XDP socket is created with the Linux socket() syscall
- Multiple sockets can process packets in parallel, each one assigned to a NIC queue
- The sockets can be passed to the eBPF program in a *BPF_MAP_TYPE_XSKMAP*



► Memory model (simplified)




UMEM is allocated by user space application with a malloc.
It has 2 rings to handle its chunks ownership between kernel and userspace:
FILL and COMPLETION rings.



Actions

eBPF program return code:

- XDP_PASS (handled by OS net stack)
- XDP_DROP (dropped)
- XDP_TX (edited and retransmitted)
- XDP_ABORTED (eBPF program error code)
- XDP_REDIRECT (forwarded to another interface/socket)



► eBPF & XDP for Windows

- Implementation is different compared to Linux
- Limited features
- Currently they aren't part of Windows releases, they must be installed by the user

microsoft/ebpf-for-windows

eBPF implementation that runs on top of Windows

52
Contributors

251
Issues

71
Discussions

3k
Stars

254
Forks



microsoft/xdp-for-windows

XDP speeds up networking on Windows

20
Contributors

1
Used by


3
Discussions

415
Stars

53
Forks



Ethernet



► Parser & Serializer

```
if len(receivedData) < 14 {
    return errors.New("invalid ethernet
frame")
}

destinationAddress := receivedData[:6]
sourceAddress := receivedData[6:12]
etherType := binary.BigEndian.Uint16(
    receivedData[12:14],
)
payload := receivedData[14:]

dataToSend := make(
    []byte, 14+len(payload),
)
copy(dataToSend[:6], destinationAddress)
copy(dataToSend[6:12], sourceAddress)
binary.BigEndian.PutUint16(
    dataToSend[12:14], uint16(0x0800),
)
copy(dataToSend[14:], payload)
```

Bypassing the Linux net stack with Go

“SouPPP” project

Domestic proxies project



▶ PPPoE Point-to-Point Protocol over Ethernet


Used by ISP to establish a connection with the home modem

Discovery (Ethernet Type 0x8863)

Initial handshake: Client sends an Ethernet broadcast message, Server sends a reply to the Ethernet source address.

Session (Ethernet Type 0x8864)

Every IP packet is encapsulated in a PPP frame and sent/received over Ethernet with this ethernet type.



► Features

- Multiple PPPoE sessions to obtain multiple IP addresses
- PPP session is handled directly by the user space application, including the connection control frames

► XDP program (C, eBPF bytecode)

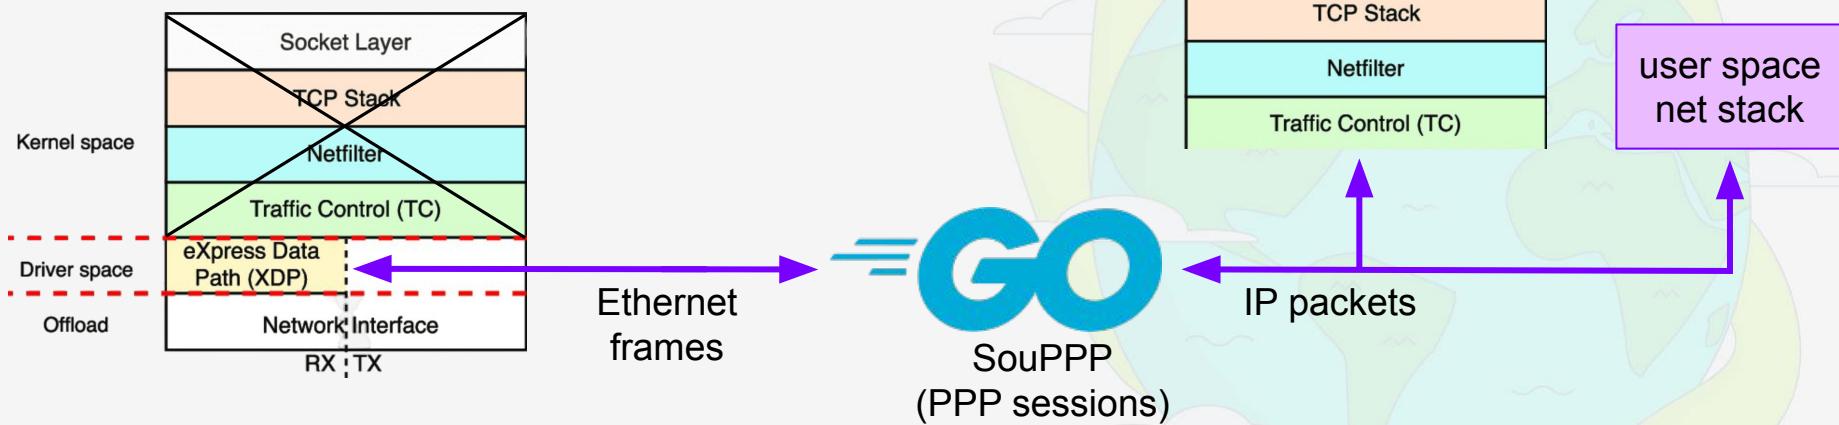
When an Ethernet frame is received:

- Redirect it to the target AF_XDP socket if Ethernet Type is PPPoE Discovery or Session
- Use XDP_PASS for all the other Ethernet Types, so the kernel can handle it “normally”

► Network traffic

Two alternatives:

- User space network stack (e.g. gVisor netstack)
- TUN/TAP virtual interface



► HTTP requests

```
dialer := &net.Dialer{ // We use it in &http.Transport{} passed to a &http.Client{}  
    // Called after creating network connection, but before dialing  
    ControlContext: func(ctx context.Context, network, address string, c  
    syscall.RawConn) error {  
        var err error  
        // Linux specific syscall to bind Dial() call to a specified network interface  
        if ctrlErr := c.Control(func(fileDescriptor uintptr) {  
            err = syscall.BindToDevice(int(fileDescriptor), "tunTapDeviceName")  
        }); ctrlErr != nil {  
            return ctrlErr  
        }  
        return err  
    }  
}
```

Bypassing the Linux net stack with Go

XDP in Go



Erik Pellizzon

erikpellizzon@gmail.com

[LinkedIn](#)




Freelancer


Looking for
collaborations



eBPF program



Code structure



► Initialization

```
interfaceData, err := netlink.LinkByName("eth0") // github.com/vishvananda/netlink
interfaceIndex := interfaceData.Attrs().Index

ebpfProg, err := func() (*xdp.Program, error) {
    spec, err := ebpf.LoadCollectionSpecFromReader(ebpfBytecodeReader)
    col, err := ebpf.NewCollection(spec)
    err = colMaps["etherTypeInputMap"].Put(0x8863, 1) // PPPoE Discovery
    err = colMaps["etherTypeInputMap"].Put(0x8864, 1) // PPPoE Session
    return &xdp.Program{Program: col.Programs[...], Queues: colMaps[...], Sockets:
        col.Sockets[...]}, nil
}()

// Attach the XDP Program to the network interface
err = ebpfProg.Attach(interfaceIndex)
```

Sockets

```
queuesNum, err := getInterfaceQueuesNumber("eth0")
if err != nil {
    return err
}
queueIDList := []int{0, 1} // When queuesNum is 2
for _, queueID := range queueIDList {
    socket, err := xdp.NewSocket(interfaceIndex int, queueID int, options
*xdp.SocketOptions)
    if err != nil {
        return err
    }
    if err := ebpfProg.Register(queueID, socket.FD()); err != nil {
        return err
    }
}
```

▶ Transmission

```
for { // One goroutine for each socket
    data := <-framesToSendChan; numFrames := 1
    descriptors := socket.GetDescs(numFrames, false) // n int, rx bool
    if len(descriptors) < numFrames { return }
    // GetFrame() returns the underlying []byte buffer of the descriptor
    copy(socket.GetFrame(descriptors[0]), data)
    descriptors[0].Len = uint32(len(data))
    numSubmitted := socket.Transmit(descriptors)
    if numSubmitted != numFrames { return }
    numTx, err := polling(socket, -1, false) // Wait for unix.POLLOUT event
    if err != nil { return }
}
```

► Receive

```
for { // One goroutine for each socket
    if n := socket.NumFreeFillSlots(); n > 0 {
        descriptors := socket.GetDescs(n, true)
        socket.Fill(descriptors) // Give ownership to the kernel
    }
    numRx, err := polling(socket, -1, true) // Wait for unix.POLLIN event
    if err != nil { return }
    if numRx <= 0 { continue }
    rxDescriptors := socket.Receive(numRx)
    for i := 0; i < len(rxDescriptors); i++ {
        packetData := slices.Clone(socket.GetFrame(rxDescriptors[i]))
        receivedFramesChan <- packetData
    }
}
```

Polling

```
func polling(socket *xdp.Socket, timeout int, rx bool) (int, error) {
    var events int16
    if rx && socket.NumFilled() > 0 { events |= unix.POLLIN }
    if !rx && socket.NumTransmitted() > 0 { events |= unix.POLLOUT }
    if events == 0 { return 0, nil }
    pollingFileDescriptor := []unix.PollFd{{Fd: int32(socket.FD()), Events: events}}
    var err error = unix.EINTR
    for errors.Is(err, unix.EINTR) { _, err = unix.Poll(pollingFileDescriptor, timeout) }
    if err != nil { return 0, err }
    if rx { return socket.NumReceived(), nil }
    numCompleted := socket.NumCompleted()
    if numCompleted > 0 { socket.Complete(numCompleted) }
    return numCompleted, nil
}
```

Greetings and conclusion



Erik Pellizzon

erikpellizzon@gmail.com

[LinkedIn](#)



Freelancer

Looking for
collaborations

